

Practical 1: Regression

Predicting the Efficiency of Organic Photovoltaics

Antonio Copete (acopete@cfa.harvard.edu, Camelot.ai Copete)
Fangli Geng (fgeng@g.harvard.edu, Camelot.ai fangligeng)
Xihan Zhang (xihanzhang@hsph.harvard.edu, Camelot.ai Xihan)

February 10, 2018

1 Technical Approach

We began by performing exploratory analysis on the sample dataset we were given, which consisted of a training set of 1 million molecules with 256 binary predictors, in addition to their SMILES string and the HOMO–LUMO gap we were seeking to predict. The test set consisted of 824,230 molecules with the same set of predictors as well as their SMILES string. The sample code we were given implemented a default Linear Regression model on the full set of 256 predictors, yielding an $R^2_{\text{LR}} = 0.461$ ($\text{RMSE}_{\text{LR}} = 0.298$) over the full training set, and a default Random Forest regression with $R^2_{\text{RF}} = 0.554$ ($\text{RMSE}_{\text{RF}} = 0.272$).

Initial inspection found that out of 256 molecular features, 221 of them were unexpressed (i.e. had $x_i = 0$) for *all* molecules, both in the training set and the test set. Dropping unimportant features would normally call for K -fold cross-validation across the training set to ensure those features are consistently unimportant in all cases. However, in the case of null values of certain features for every element of the training set, it is not only legitimate but necessary to drop those features from all further analysis, as fitting along null dimensions would constitute a form of overfitting.

Reducing sample dataset to 31 expressed features, we performed both regularized and non-regularized linear regression with cross-validation¹ and hyperparameter tuning, under following methods:

1. Non-regularized linear regression: Yields a mean $R^2 = 0.461$ from 5-fold cross-validation, similar to the original result with 256 predictors.
2. Ridge Regression: Tuning the hyperparameter α by grid-search cross-validation results in an optimal value of 5.86×10^{-5} , with a best $R^2 = 0.459$, a result almost identical to non-regularized regression.
3. Elastic Net: Trying a more general regularization of the form $\alpha_1 L_1$ (Lasso) + $\alpha_2 L_2$ (Ridge regression), and tuning the L_1 ratio $\equiv \alpha_1/\alpha_2$ by grid-search cross validation, yields an optimal value of 0.0, which indicates that that a regularization biased towards sparse models, such as the Lasso, might not be the best choice in this case.

The results from linear regression on the sample dataset consistently gave us low predictive value, even when compared to generic random forest regression, which led us into 2 parallel tracks:

1. Feature engineering:

The RDKit package in Python use the SMILES string to provide an array of information and functions about a molecule, including substructure searching, chemical reaction, enumeration of molecular resonance structures, and several sorts of fingerprints. From our research we decided to extract the *Morgan fingerprint function*, which is widely used with SKLearn in ML and carries the information we considered to be most relevant to the efficiency of organic photovoltaics, such as: (a) Connectivity: Element, numbers of heavy neighbors, numbers of Hs, charge, isotope, inRing;

¹After initially trying 3-fold, 5-fold and 10-fold cross-validation, we settled on 5-fold cross-validation as the best compromise between accuracy and computational speed.

(b) Chemical features: Donor, Acceptor, Aromatic, Halogen, Basic, Acidic; and (c) Taking into account the neighborhood of each atom.

The process takes the “SMILES” column from the training and test data and uses RDKit to convert it into molecule objects, followed by gathering of the Morgan fingerprint by use of the command `rdkit.Chem.getmorganfingerprintasbitvect()`, adjusting the parameter `nBits` to 512, 2048 and 4096 aiming to test datasets with various precision levels. After iterating this process for each string, we appended the columns together to form a new input matrix. The size of the dataset greatly varies with precision level, and for this reason we had to slice the datasets into 8 parts in order to accommodate computational constraints.

2. Non-linear methods (among these we concentrated on 2 broad categories):

(a) Tree-based Ensemble methods:

Binary predictors are well suited to the branch-like structure of decision trees. Since deep may overfit the training set, we applied ensemble methods to average over the predictions of all trees. The baseline ensemble methods are: (a) Bagging: bootstrap multiple training sets to fit multiple trees and reduce the variance. (b) Random Forest: further reduce the variance by randomly picking a subset of \sqrt{p} predictors from the total p at each split. (c) Extremely Randomized Trees: further reduce the variance by selecting cutting points from the subsets totally at random. (d) AdaBoost: use a linear combination of multiple simple trees to reduce the residual step by step.

Baseline models were fitted on the 31 expressed features. A 20% validation set was used to test the MSE generated by each model. Bagging, Random Forest, Extremely Randomized Trees are quite similar to each other. AdaBoost, which focuses on reducing bias rather than variance, is much weaker. Thus, we selected Random Forest as main model, and tuned the number of estimators (i.e. the total number of trees to average over) to reduce the variance of the model. 5-fold cross validation was used on 1/8 of the total sample with 2048 newly generated features.

From the results, we concluded that 13 estimators is sufficient to make the MSE of the test set close to the training set. Given that the total sample size is 8 times that of the test set where we tuned the model, setting the number of estimators to 150 for the whole data set is reasonable.

The distribution of predictor importances from the baseline random forest is plotted below:




Figure 1: Distribution of feature importances from Random Forest

Though most predictors have little importance (10^{-8} – 10^{-2}), combining them can result in good precision on the validation set ($R^2 = 0.936$, RMSE = 0.1029). This made us want to explore further into Deep Learning methods, which consider interactions among features.

(b) Deep learning:

In order to model the interactions that would be expected among the broad set of molecular properties extracted from RDKit, we decided to explore the training of simple neural networks and got acquainted with the basics of Deep Learning as a tool for this problem. In all cases, we used an architecture of one or more dense (fully-connected) hidden layers between the input layer of N parameters and the output layer of 1 node for the response we wanted to model. For each node in the hidden layers, we used the widely used “relu” function $f(x) = \begin{cases} 0 & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases}$ as activation function. We trained the model using the *keras* module in Python, using the *Adam* optimizer for the learning rate, setting aside 20% of training data for validation. The model fits for the weights of the connections between pairs of nodes in the network, which in the case of 1 hidden layer with M nodes, would result in $(N + 1) \times M$ parameters to be trained.

Taking a trial-and-error approach to the network architecture, we show the results of validation set MSE for 3 broad cases, using the 31 expressed features in the sample training dataset: 1) 1 hidden layer with variable number of nodes, 2) 2 hidden layers, with 31 nodes on the first and variable number of nodes on the second, and 3) 3 hidden layers, with 31 nodes on each of the first two, and variable number of nodes on the third.

The result of this experiment showed that the MSE was not as sensitive to the addition of a second and third layer, as it was to the number of nodes in the first layer of the 1-hidden layer architecture. Given the computational limitations that were expected for a much larger set of features, we used this result to use the feature-engineered set to train only 1-layer models with a limited number of nodes in the same range as the number of predictors.




Figure 2: MSEs of Neural Network architectures

2 Results

Figure 3 shows the overall results we obtained from a wide array of regression methods on different sets of features, as described in the Technical Approach section, and roughly listed in the order in which we tried them.

One evident result is the marked improvement in the predictive value of feature-engineered models employing a large number of predictors (429/1571/2048/4096), compared to those from the sample dataset (31/256). Within models of comparable numbers of predictors, linear regression performed relatively poorly, indicating the non-linear interactions between features that are likely to result in the correct HOMO–LUMO gaps. Our 2 best results in the Camelot table came from employing neural network models, one with 4096 features and 1 hidden layer of 1000 nodes, and another with 1571 features and 1 hidden layer of 500 nodes. These results ranked 9 (RMSE = 0.03461) and 11 (RMSE = 0.0359) in the final table, both in the top 10% of scores.

3 Discussion

We started with exploratory analysis on the sample dataset with 256 features by fitting Linear Regression and Random forest as baseline models, and tried to understand the basic relation-




Figure 3: Overall results for different regression methods and number of predictors

ships between the data and given features. Then we started tuning the parameters within linear regressions and Random Forest to see how much we could improve upon the current data. After seeing the insignificant improvements, we started to experimenting in Deep Learning model with selected important features. We found Deep Learning model provided the most accurate results so far, and the reason would be allowing features interactions within molecules. Therefore, we decided to focus on Neural Network models further on. Meanwhile, we pulled Morgan Fingerprint features by RDKit in two sizes of Bits and produced 2048 and 4096 features.

Considering computational power and time, we also tried to drop the features within 2048 features, based on feature importance yield from Regression RF models. Finally, based on the all previous experiments, we set up 3 configurations for neural network models: 2 layer with 2048 features and 100 units in each layer (RMSE=0.03564), 1 layer with 2048 features with 1000 units (RMSE=0.03553), 1 layer with 4096 features with 1024 units (RMSE=0.03461). Comparing the former two, we found the layer number is still not as important as unit number, which is consistent with what we found using sample dataset. Comparing the latter two, we found 2048 Morgan fingerprint is basically informative enough to predict, and 4096 features in more detail couldn’t improve the prediction as much but meanwhile taking a much larger amount of time for computation.

Additionally, we turned to find additional features and found counted-based fingerprints could be a good option in addition to the bit-based fingerprint, since it count-based fingerprint count the number of times a feature appears instead of simply that it appears. It may provide new and important information to our prediction. Furthermore, we wanted to try other algorithms like Principal Component Analysis to select more important features to reduce computational complexity, and to implement neural network models with 2 layer and 2000 nodes to allow more complex interactions between selected important features. However, due to the time constraints we had, we didn’t realize those plans.

As a summary, we found that the linear regression and random forest are good methods to rapidly select useful features. Neural network is much more accurate in non-linear regression, especially with large dataset and rich features condition.